

Linux Device Drivers

An Introduction

Team Emertxe



Introduction



Familiarity Check



- Good C & Programming Skills
- Linux & the Filesystem
 - Root, User Space Headers & Libraries
- Files
 - Regular, Special, Device
- Toolchain
 - gcc & friends
- Make & Makefiles
- Kernel Sources (Location & Building)

The Flow

- Introduction
 - Linux kernel Ecosystem
 - Kernel Source Organization
 - Command set and Files
 - Writing the first driver (Module)
- Character Drivers
 - Device files
 - Device access from user space (End to End flow)
 - Registering the driver
 - File Operations and registration
 - Data transfer between User and Kernel space
 - ioctl
- Memory & Hardware
- Time & Timings
- USB Drivers
- Interrupt Handling
- Block Drivers
- PCI Drivers
- Debugging

The Flow...

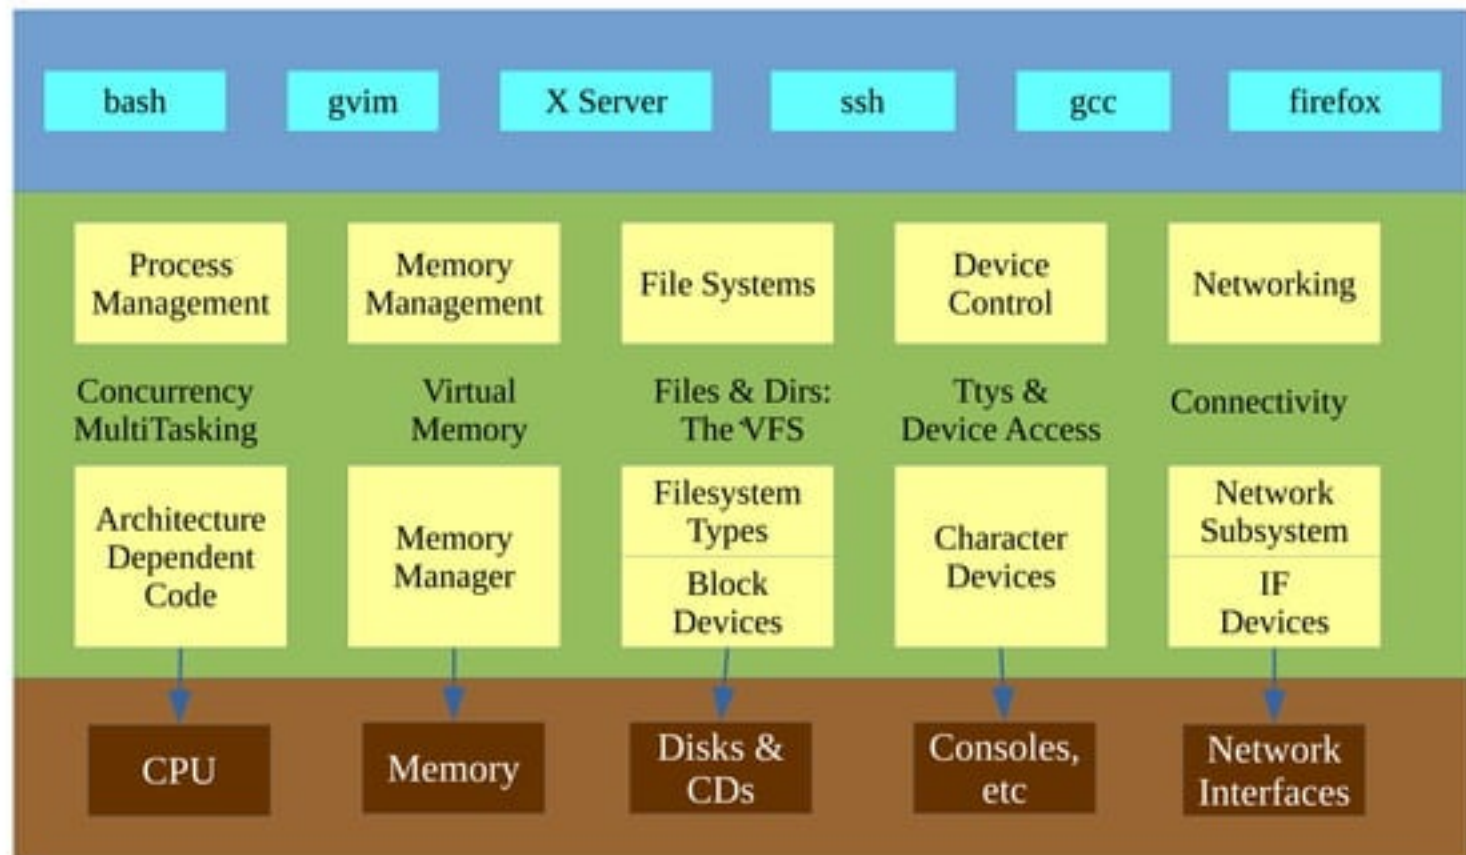
- Introduction
 - Linux kernel Ecosystem
 - Kernel Source Organization
 - Command set and Files
 - Writing the first driver (Module)
- Character Drivers
 - Device files
 - Device access from user space (End to End flow)
- Memory & Hardware
- Time & Timings
- USB Drivers
- Interrupt Handling
- Block Drivers
- PCI Drivers
- Debugging

Hands-On

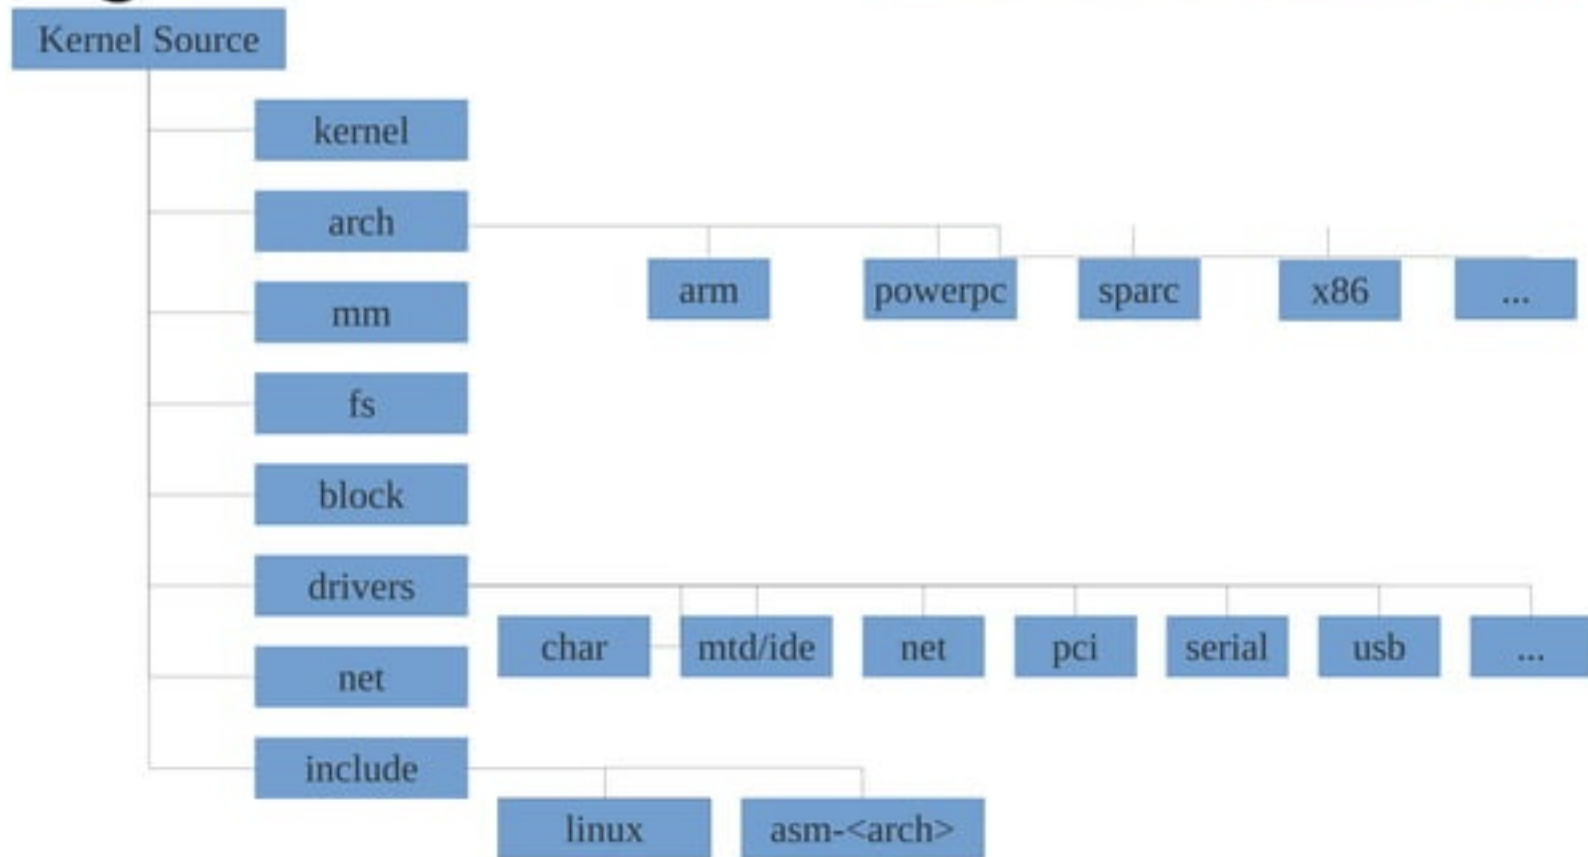


- Your First Driver
- Character Drivers
 - Null Driver
 - Memory Driver
 - UART Driver for Customized Hardware
- USB Drivers
 - USB Device Hot-plug-ability
 - USB to Serial Hardware Driver
- Filesystem Modules
 - VFS Interfacing
 - “Pseudo” File System with Memory Files

Linux Driver Ecosystem



Kernel Source Organization



The Locations & Config Files



- Kernel Source Path: `/usr/src/linux`
- Std Modules Path:
 - `/lib/modules/<kernel version>/kernel/...`
- Module Configuration: `/etc/modprobe.conf`
- Kernel Windows:
 - `/proc`
 - `/sys`
- System Logs: `/var/log/messages`

The Commands



- lsmod
- insmod
- modprobe
- rmmod
- dmesg
- objdump
- nm
- cat /proc/<file>

The Kernel's C



- ctor & dtor
 - init_module, cleanup_module
- printf
 - printk
- Libraries
 - <kernel src>/kernel
- Headers
 - <kernel src>/include

The Init Code

```
static int __init mfd_init(void)
{
    printk(KERN_INFO "mfd registered");
    ...
    return 0;
}
module_init(mfd_init);
```

The Cleanup Code

```
static void __exit mfd_exit(void)
{
    printk(KERN_INFO "mfd deregistered");
    ...
}
module_exit(mfd_exit);
```

Usage of printk

- <linux/kernel.h>
- Constant String for Log Level
 - KERN_EMERG "<0>" /* system is unusable */
 - KERN_ALERT "<1>" /* action must be taken immediately */
 - KERN_CRIT "<2>" /* critical conditions */
 - KERN_ERR "<3>" /* error conditions */
 - KERN_WARNING "<4>" /* warning conditions */
 - KERN_NOTICE "<5>" /* normal but significant condition */
 - KERN_INFO "<6>" /* informational */
 - KERN_DEBUG "<7>" /* debug-level messages */
- printf like arguments

The Other Basics & Ornaments



- Headers
 - `#include <linux/module.h>`
 - `#include <linux/version.h>`
 - `#include <linux/kernel.h>`
- `MODULE_LICENSE("GPL");`
- `MODULE_AUTHOR("Emertxe");`
- `MODULE_DESCRIPTION("First Device Driver");`

Building the Module



- Our driver needs
 - The Kernel Headers for Prototypes
 - The Kernel Functions for Functionality
 - The Kernel Build System & the Makefile for Building
- Two options
 - Building under Kernel Source Tree
 - Put our driver under drivers folder
 - Edit Kconfig(s) & Makefile to include our driver
 - Create our own Makefile to do the right invocation

Our Makefile

```
ifneq (${KERNELRELEASE},)
    obj-m += <module>.o
else
    KERNEL_SOURCE := <kernel source directory path>
    PWD := $(shell pwd)
default:
    $(MAKE) -C ${KERNEL_SOURCE} SUBDIRS=$(PWD) modules
clean:
    $(MAKE) -C ${KERNEL_SOURCE} SUBDIRS=$(PWD) clean
endif
```

Try Out your First Driver

Character Drivers



Major & Minor Number



- `ls -l /dev`
- Major is to Driver; Minor is to Device
- `<linux/types.h>` ($\geq 2.6.0$)
 - `dev_t`: 12 & 20 bits for major & minor
- `<linux/kdev_t.h>`
 - `MAJOR(dev_t dev)`
 - `MINOR(dev_t dev)`
 - `MKDEV(int major, int minor)`

Registering & Unregistering



- Registering the Device Driver
 - `int register_chrdev_region(dev_t first, unsigned int count, char *name);`
 - `int alloc_chrdev_region(dev_t *dev, unsigned int firstminor, unsigned int cnt, char *name);`
- Unregistering the Device Driver
 - `void unregister_chrdev_region(dev_t first, unsigned int count);`
- Header: `<linux/fs.h>`

The file operations

- `#include <linux/fs.h>`
- `struct file_operations`
 - `int (*open)(struct inode *, struct file *);`
 - `int (*release)(struct inode *, struct file *);`
 - `ssize_t (*read)(struct file *, char __user *, size_t, loff_t *);`
 - `ssize_t (*write)(struct file *, const char __user *, size_t, loff_t *);`
 - `struct module owner = THIS_MODULE; / linux/module.h> */`
 - `loff_t (*llseek)(struct file *, loff_t, int);`
 - `int (*ioctl)(struct inode *, struct file *, unsigned int, unsigned long);`

User level I/O



- `int open(const char *path, int oflag, ...)`
- `int close(int fd);`
- `ssize_t write(int fd, const void *buf, size_t nbyte)`
- `ssize_t read(int fd, void *buf, size_t nbyte)`
- `int ioctl(int d, int request, ...)`
 - The `ioctl()` function manipulates the underlying device parameters of special files.
 - The argument `d` must be an open file descriptor.
 - The second argument is a device-dependent request code.

The file & inode structures



- struct file
 - mode_t f_mode
 - loff_t f_pos
 - unsigned int f_flags
 - struct file_operations *f_op
 - void * private_data
- struct inode
 - unsigned int iminor(struct inode *);
 - unsigned int imajor(struct inode *);

Registering the file operations



- `#include <linux/cdev.h>`
- 1st way initialization:
 - `struct cdev *my_cdev = cdev_alloc();`
 - `my_cdev->owner = THIS_MODULE;`
 - `my_cdev->ops = &my_fops;`
- 2nd way initialization:
 - `struct cdev my_cdev;`
 - `cdev_init(&my_cdev, &my_fops);`
 - `my_cdev.owner = THIS_MODULE;`
 - `my_cdev.ops = &my_fops;`

Registering the file operations...



- The Registration
 - `int cdev_add(struct cdev *cdev, dev_t num, unsigned int count);`
- The Unregistration
 - `void cdev_del(struct cdev *cdev);`

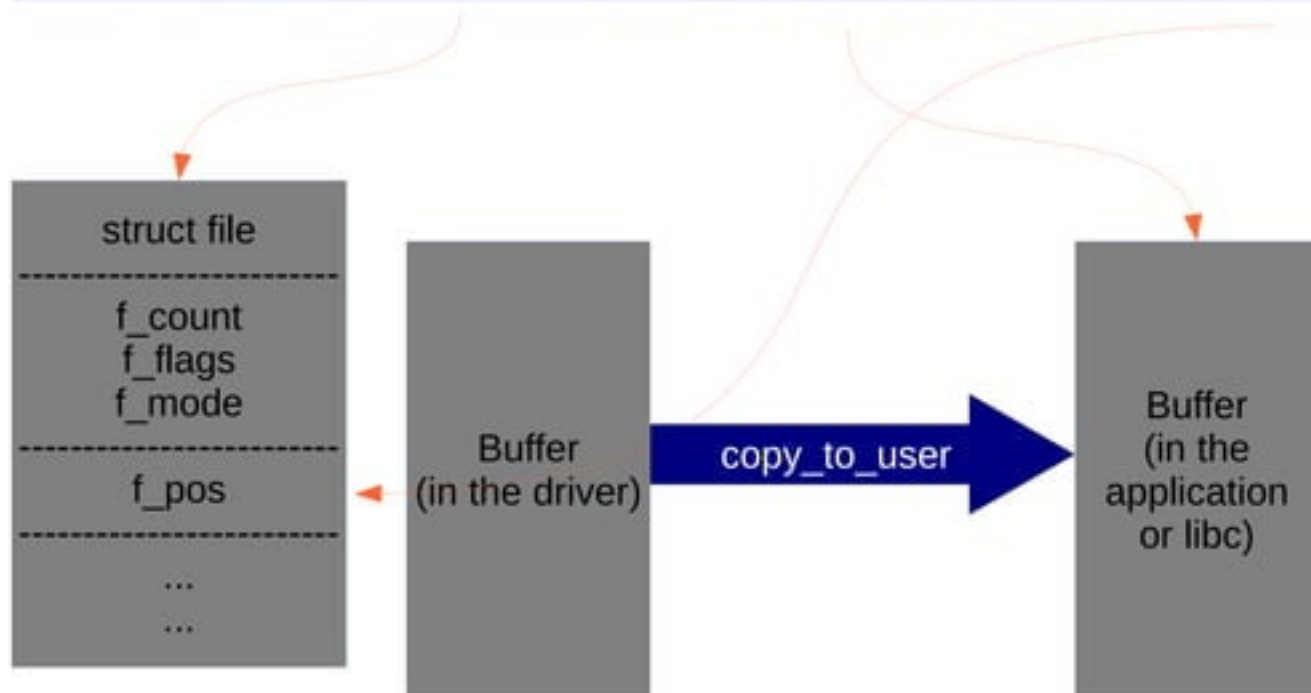
Registering/Unregistering Old Way



- Registering the Device Driver
 - `int register_chrdev(undigned int major, const char *name, struct file_operations *fops);`
- Unregistering the Device Driver
 - `int unregister_chrdev(undigned int major, const char *name);`

The read flow

```
ssize_t my_read(struct file *f, char __user *buf, size_t cnt, loff_t *off)
```



Kernel Space (Non-swappable)

User Space (Swappable)

The /dev/null read & write

```
ssize_t my_read(struct file *f, char __user *buf, size_t cnt, loff_t
    *off)
{
    ...
    return read_cnt;
}

ssize_t my_write(struct file *f, char __user *buf, size_t cnt, loff_t
    *off)
{
    ...
    return wrote_cnt;
}
```

The mem device read

```
ssize_t my_read(struct file *f, char __user *buf, size_t cnt, loff_t
    *off)
{
    ...
    if (copy_to_user(buf, from, cnt) != 0)
    {
        return -EFAULT;
    }
    ...
    return read_cnt;
}
```

The mem device write

```
ssize_t my_write(struct file *f, char __user *buf, size_t cnt, loff_t
    *off)
{
    ...
    if (copy_from_user(to, buf, cnt) != 0)
    {
        return -EFAULT;
    }
    ...
    return wrote_cnt;
}
```

Dynamic Device Node & Classes



- Class Operations
 - `struct class *class_create(struct module *owner, char *name);`
 - `void class_destroy(struct class *cl);`
- Device into & Out of Class
 - `struct class_device *device_create(struct class *cl, NULL, dev_t devnum, NULL, const char *fmt, ...);`
 - `void device_destroy(struct class *cl, dev_t devnum);`

The I/O Control API



- `int (*ioctl)(struct inode *, struct file *, unsigned int cmd, unsigned long arg)`
- `int (*unlocked_ioctl)(struct file *, unsigned int cmd, unsigned long arg)`
- Command
 - `<linux/ioctl.h> -> ... -> <asm-generic/ioctl.h>`
 - Macros
 - `_IO, _IOR, _IOW, _IOWR`
 - Parameters
 - type (Magic character) [15:8]
 - number (index) [7:0]
 - size (param type) [29:16]

The I/O Control API

- Macro Usage

`_IO(type, index)`


`[_IOR | _IOW | _IOWR](type, index,
datatype/size)`

Module Parameters



- `<linux/moduleparam.h>`
 - Macros
 - `module_param(name, type, perm)`
 - `module_param_array(name, type, num, perm)`
 - Perm (is a bitmask)
 - 0
 - `S_IRUGO`
 - `S_IWUSR | S_IRUGO`
 - Loading
 - `insmod driver.ko name=10`

x86 Architecture




Memory Access



Physical Vs Virtual Memory



- The kernel Organizes Physical memory in to pages
 - Page size Depends on Arch
 - X86-based 4096 bytes
- On 32-bit X86 system Kernel total Virtual address space
 - Total 4GB (pointer size)
 - Kernel Configuration Splits 4GB in to
 - 3GB Virtual Sp for US
 - 1GB Virtual Sp for Kernel
 - 128MB KDS
 - Virtual Address also called



Memory Access from Kernel Space



- Virtual Address on Physical Address
 - #include <linux/gfp.h>
 - unsigned long __get_free_pages(flags, order); etc
 - void free_pages(addr, order); etc
 - #include <linux/slab.h>
 - void *kmalloc(size_t size, gfp_t flags);
 - GFP_ATOMIC, GFP_KERNEL, GFP_DMA
 - void kfree(void *obj);
 - #include <linux/vmalloc.h>
 - void *vmalloc(unsigned long size);
 - void vfree(void *addr);

Memory Access from Kernel Space...



- Virtual Address for Bus/IO Address
 - #include <asm/io.h>
 - void *ioremap(unsigned long offset, unsigned long size);
 - void iounmap(void *addr);
- I/O Memory Access
 - #include <asm/io.h>
 - unsigned int ioread[8|16|32](void *addr);
 - unsigned int iowrite[8|16|32](u[8|16|32] value, void *addr);
- Barriers
 - #include <linux/kernel.h>: void barrier(void);
 - #include <asm/system.h>: void [r|w|]mb(void);

Hardware Access

A large, horizontal arrow graphic pointing to the right. The arrow has a gradient from magenta on the left to dark purple on the right. It features a white outline and a smaller, offset arrow shape within it, creating a sense of depth and movement.

I/O Accesses from Kernel Space



- I/O Port Access
 - `#include <asm/io.h>`
 - `unsigned in[b|w|l](unsigned port);`
 - `void out[b|w|l](unsigned [char|short|int] value, unsigned port);`

Hands-On the Hardware

